

# ECE344: Operating Systems

Arnav Patil

University of Toronto

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>The Kernel</b>	<b>2</b>
<b>3</b>	<b>Libraries</b>	<b>3</b>
<b>4</b>	<b>Process Creation</b>	<b>3</b>
<b>5</b>	<b>Process Management</b>	<b>4</b>
<b>6</b>	<b>Basic IPC</b>	<b>5</b>
<b>7</b>	<b>Process Practice</b>	<b>6</b>
<b>8</b>	<b>Subprocesses</b>	<b>7</b>
<b>9</b>	<b>Basic Scheduling</b>	<b>8</b>
<b>10</b>	<b>Advanced Scheduling</b>	<b>9</b>
<b>11</b>	<b>Virtual Memory</b>	<b>11</b>
<b>12</b>	<b>Page Tables</b>	<b>12</b>
<b>13</b>	<b>Page Table Implementation</b>	<b>13</b>
<b>14</b>	<b>Virtual Memory Lab Primer</b>	<b>14</b>
<b>15</b>	<b>Priority Scheduling and Memory Mapping</b>	<b>15</b>

# 1 Introduction

- Any software that gets written either:
  - is the operating system, or
  - interacts with the operating system.
- The OS as a resource manager:
  - allows multiple programs to execute at the same time
  - manages/protects memory, IO devices, etc

## Three Core OS Concepts

- Virtualization – share a memory/resource by mimicking multiple independent copies
- Concurrency – handles multiple things happening
- Persistence – retains data consistency w/o power

# 2 The Kernel

- The OS provides the illusion that each program has full access to all resources – on its own machine
- Called virtualization because one physical machine but illusion of multiple virtual machines
- Kernel mode is a privilege level on a CPU that gives access to more instructions
  - Different architectures call it different things
  - Instructions allow only trusted SW to interact with the HW
- Program – file containing instructions and data
- Process – instance of a process being executed
- User mode > supervisor > hypervisor > machine

## System Calls

- Transitions b/w user and kernel mode, OS API
  - Create/destroy threads, allocate/deallocate memory, etc.
- API – abstracts details and describes arguments and return value of a function
- ABI – specifies details, specifically how to pass arguments and where the return value is
- E.g. say a program calls `read()`
  - Execution goes via library and issues ‘heap’
  - Trap invokes the kernel which accesses the disk
  - Kernel returns the results to the program
- The kernel is a long, constantly running program
  - Link using libraries, there’s no `main()`
  - Lets you load code (modules)
  - Code executes on-demand

- If you write a kernel module, you can execute privileged instructions
- Monolithic kernels have less features in user mode, microkernel kernels have more
- Hybrid kernels are between monolithic and microkernel
- Nano/pico kernels have even more features in user mode
- ISAs – x86/64 (amd64), arm64 (aarch64), riscv

## File Descriptors

- IPC – inter-process communication is transferring data between processes
- File descriptor – resource that uses read/write (stores as an index)

## 3 Libraries

- Libraries are used as part of the OS
- Apps may pass through multiple layers of libraries
- An OS consists of the kernel and libraries required for your applications

## Dynamic Libraries Are For Reusable Code

- C standard library is a dynamic library (`.so`) like any other on the system
  - Collection of `.o` files containing function definitions
- OS loads `libc.so` in memory only once during boot

## Static vs Dynamic Libraries

- Drawbacks to static:
  - Statically linking prevents re-using libraries
  - Any updates to a static library requires the executable to be recompiled

## Dynamic Library Updates Can Break Executables

- A dynamic library update may subtly break an ABI causing a crash
- `structs` are laid out in memory w/ the fields matching the order of declaration
- Semantic versioning meets developer's expectations
  - Given a version number MAJOR.MINOR.PATCH increment:
    - \* MAJOR when you make incompatible API/ABI changes
    - \* MINOR when you add functionality in a backwards compatible manner
    - \* PATCH when you make backwards compatible bug fixes

## 4 Process Creation

- Process is a running instance of a program
  - Virtual registers
  - Virtual memory
  - File descriptors – an array of numbers that points to files that the kernel is managing

## Process Control Blocks

- In Linux this is the `task_struct`
- Contains process state, CPU registers, scheduling, memory management, IO status info
- Each process gets a unique PID assigned to it
- We can read processes using the `proc` filesystem
  - `/proc` doesn't contain real files, but we can use it as such
  - Every directory that's a number is a currently running process (PID)
- Windows – we load program into memory and create the PCB
- Linux – decomposes process creation into more flexible abstractions

## Cloning a Process

- Pause currently running process, copy its PCB into a new one. This reuses all info from the old process, including variables
- Distinguished by a parent-child relationship
- `int fork(void)` creates a new process. Returns:
  - -1 on failure
  - 0 in the child process
  - the child's PID in the parent process
- `execve()` replaces a process with another program
  - `pathname` – path of the program to load
  - `argv` – array of strings, arguments to process
  - `envp` – same as `argv` but for environment
- Modern OS's are smart and won't let you make infinite forks.

# 5 Process Management

## Linux Process Management

- Can read process by doing `/proc/<PID>/states/` `grep state`
  - R – running and runnable
  - S – interruptible sleep
  - D – uninterruptible sleep
  - T – stopped
  - Z – zombie
- The kernel allows us to explicitly stop processes but we must restart them
- After the kernel initializes, it creates a single process
  - Looks for program in `/sbin/init`
- `init` is responsible for executing every other process, it must always be active, else kernel thinks you're shutting down
- Using `htop` helps us keep track of existing processes
- Kernels will eventually recycle PIDs of finished processes

## Maintaining the Parent-Child Relationship

- OS sets exit status to a process that's finished executing
- Minimum acknowledgment the parent has to do is read the child's exit status
- Two possible situations:
  - Child exits first (zombie process)
  - Parent exits first (orphan process)
- `wait(status)` – where to store wait status of the process
  - -1 on failure
  - 0 for non-blocking calls w/ no child change
  - PID of child process with a change
- OS can only remove the zombie's entry after `wait()` retrieves its PID

## Zombie Processes

- Process has been terminated but hasn't been acknowledged
- Process may have an error where it never reads the child's exit status
- OS can interrupt the parent process to acknowledge the child
- This is a suggestion, and the parent is free to ignore it
  - Basic form of IPC, called a signal
- The OS has to keep a zombie process until it's been acknowledged
- If the parent process ignores it, the zombie has to wait to be reparented

## Orphan Processes Need a New Parent

- Child processes still need a process to acknowledge its exit
- OS re-parents the child to `init` – `init` now responsible for acknowledging the child
- `init` accepts all orphans, dead or alive

## 6 Basic IPC

- IPC is transferring bytes between two or more systems
- Reading/writing is a form of IPC
- `read` just reads data from a file descriptor
- No EOF character, just returns 0 bytes read
  - Kernel returns 0 on a closed file descriptor
- We need to check for errors
- `write` similarly returns the number of bytes written, but we can't always assume success

## Standard File Descriptors

- We could close fd 0 (standard input) and open a file instead
- Signals are a form of IPC that interrupts
- Kernel sends a number to your program indicating the type of signal
  - Kernel's handlers either ignore the signal or terminate
  - If the default handler occurs, the exit code will be 128 + the signal number
- Ctrl+C sends SIGINT – signal interrupt from keyboard

## Setting Own Signal Handlers

- Declare a function w/ no return and 1 int argument
- Some common interrupts on Linux:
  - 2: SIGINT (keyboard)
  - 9: SIGKILL (terminate)
  - 11: SIGSEGV (seg fault)
  - 15: SIGTERM (terminate)
- Processes can be interrupted at any point of execution, and resumes once the signal handler returns – example of **concurrency**
- `kill PID` sends SIGTERM signal but won't terminate if the process is in uninterruptible sleep
- `kill -9 PID` will kill the process no matter what

## Non-Blocking Calls

- A non-blocking call returns immediately so we can check if something happens
- To turn `wait` into a non-blocking call, we can use the flag `WNOHANG` in options
- To react to changes in a non-blocking call we can either use polling or an interrupt

## 7 Process Practice

- Uniprogramming is for old-batch OSs
- Uniprogramming is when only one process is running at a time – no parallelism and no concurrency
- Multiprocessing – parallel or concurrent both possible, we want parallel AND concurrent

## Scheduler Decides When to Switch

- To create a process, the OS has to at least load it into memory
- While maintaining, the scheduler decides when it's running
- First we focus on mechanics of switching processes

## Core Scheduling Loop

1. Pause currently running process
  2. Save its state so you can restore later
  3. Get next process to run from scheduler
  4. Load next process' state
- Cooperative multitasking is when the process uses a syscall to tell the OS to pause it
  - True multitasking is when the OS retains control and pauses processes

## Context Switching

- Name for switching processes
- We have to save all register values using the same CPU the process is already using
- Hardware support for saving state, but we may not want to save everything
- Context switching is pure overhead, we want to minimize

## A New API

- `int pipe(int pipefd[2]);`
  - Returns 0 on success and -1 on failure (sets errno)
  - Forms a one-way communication channel using two file descriptors
  - `pipefd[0]` is read and `pipefd[1]` is the write end
- Kernel-managed buffer, any data written to one is read on the other end

## 8 Subprocesses

### We Want to Send/Receive Data From a Process

1. Create a new process that launches command line argument
2. Send string "Testing\n" to that process
3. Receive any data it writes to that process

### A More Convenient API – `execlp`

- Doesn't return on success, -1 on failure
- Will let you skip string arrays
- It will also search for executables using the PATH environment variable

### Final APIs – `dup` and `dup2`

- Returns a new file descriptor on success, -1 on failure and sets errno
- Copies file descriptor so both refer to the same thing

## 9 Basic Scheduling

### Preemptible and Non-Preemptible Resources

- Preemptible resources can be taken away and used for something else e.g. a CPU
- The resource is then shared through scheduling
- A non-preemptible resource can't be taken away w/o acknowledgment e.g. disk space
- The resource is instead shared through allocations and deallocations
  - Parallel and distributed systems may allow you to allocate a CPU

### Dispatchers and Schedulers Work Together

- A dispatcher is a low-level mechanism responsible for context switching
- A scheduler is a high-level policy responsible for deciding which processes to run

### Scheduler Runs Whenever a Process Changes State

- For non-preemptible resources – process runs until completion, once started
- Scheduler only makes decision once the process is terminated
- Preemptive allows the OS to run scheduler at will

### Important Metrics

- Minimize waiting time and response time
- Maximize CPU utilization
- Maximize throughput
- Fairness

### First-Come First-Serve – FCFS

- Most basic form of scheduling
- First process that arrives gets access to CPU
- Processes stored in a FIFO queue

### Shortest Job First – SJF

- Always schedule job w/ shortest burst time first
- Still assuming no preemption
- But it is not practical
  - Likely optimal at minimizing average wait time
  - Don't know burst time of each process
  - Long jobs may be starve (or never execute)

### Shortest Remaining Time First – SRTF

- Assume that minimum waiting time is 1 unit, optimize average waiting time



## Round-Robin – RR

- Haven't discussed fairness so far – trade-offs
- OS divides execution into time slices (or quotas)
- Maintain a FIFO queue of processes similar to FCFS
- Pre-empt is still running at end of the quantum and re-add to queue
- RR performance depends on quantum length and job length
  - RR has low response times and good interactivity
  - Fair allocation of CPU and low average waiting time
  - Performance depends on quantum length
    - \* Too high – becomes FCFS
    - \* Too low – too many context switches (high overhead)
- RR has poor average waiting time when jobs have similar lengths

## Scheduling Trade-Offs

- FCFS – most basic scheduling algorithm
- SJF – tweak to reduce waiting time
- SRTF – uses SJF but with preemption
- RR – optimizes fairness and response time

# 10 Advanced Scheduling

## Adding Priorities

- We may favour some processes over others
- Run high priority processes first, round-robin processes of equal priority
- On Linux, -20 is the highest priority and 19 is the lowest
- We may lead processes to starvation if there's lots of high-priority loads
- One solution is to have the OS dynamically change the priority

## Priority Inversion

- We can accidentally change priority of low to high, would depend on if a high-priority process depended on a low-priority process
- Solution is to have priority inheritance – inherit priority of the waiting process
- Idea is to separate processes that users interact with
  - Foreground processes are interactable and need good response time
  - Background processes may just need good throughput

## Using Multiple Queues

- Create different processes for foreground and background processes
  - Foreground – RR
  - Background – FCFS
- Now we have to schedule b/w queues
  - RR the queues or use a priority system
- We'll assume symmetric multiprocessing (SMP)
  - All CPUs connected to same physical memory
  - CPUs all have their own (lowest-level) caches
- One approach is to use same scheduling for all CPUs
  - Only one scheduler – adds processes while CPU available
  - Pros: good CPU utilization, fair to all processes
  - Cons: not scalable, poor cache locality
- Another is to create per-CPU schedulers
  - Assign new processes to CPUs with the lowest # of processes
  - Pros: easy to implement, scalable, good cache locality
  - Cons: load imbalance
- We can also compromise b/w global and per-CPU
  - Keep a global scheduler that can rebalance per-CPU queues
    - \* If a CPU is idle, take a process from another CPU (work stealing)
  - We may have some processes that are more sensitive to caches
  - Using processor affinity
    - \* Preference of a process to be scheduled on the same core
- Gang scheduling (co-scheduling)
  - Multiple processes may need to be scheduled simultaneously
  - Scheduler on each CPU cannot be completely independent
  - Requires a global context-switching across all CPUs

## Real-Time Scheduling

- Real-time means there are time constraints, either for a deadline or rate
- Hard and soft real-time systems

## Linux FCFS and RR Scheduling

- Use a multilevel queue scheduler for processes with the same priority
  - Also let the OS dynamically adjust the priority
  - Soft real-time processes – always schedule for the highest priority first
  - Normal processes – adjust priority based on aging

## **O(1) Scheduling Issues**

- Now kernel has to detect processes which are interactive using heuristics
- Processes that sleep a lot may be more interactive

## **Ideal Fair Scheduling (IFS) is Fairest but Impractical**

- Performs way too many context switches
- Have to constantly scan all processes at  $O(N)$

## **Completely Fair Scheduler (CFS)**

- For each runnable process, assign it to a virtual runtime
- At each scheduling point, increase virtual runtime by  $t \times \text{weight}$  (priority)
  - Virtual time monotonically increases
    - \* Scheduler selects process based on lowest virtual runtime
    - \* Compute dynamic runtime based on IFS
- Implemented on red-black tree, self-balancing BST
  - $O(\log(n))$  insert, delete, find operations

# **11 Virtual Memory**

## **Requirements of Virtual Memory**

- Multiple processes must co-exist
- Processes unaware they are sharing physical memory and cannot access each other's data, unless explicitly allowed
- Performance close to actual physical memory
- Limit amount of wasted memory – fragmentation

## **Segmentation/Segments are Coarse-Grained**


- Each segment is a variable sized – dynamically allocated
- Legacy technique that's not really used anymore
- Segments are large and hard to relocate, but leads to fragmentation
- Each segment contains a base, limit, and permissions
- MMU checks that offset is within limit, then calculates  $\text{base} + \text{offset}$  and does permission checks
  - If not, results in a seg fault

## **Memory Management Unit (MMU)**

- Maps virtual addresses to physical addresses and does permission checks
- One technique is dividing memory into fixed size pages (4096 bytes)
- A page in virtual memory is called a page, and a page in physical memory is called a frame

## Addressing

- Typically don't use all 64 bits its virtual address
- CPUs may have diff levels of virtual addresses you can use
- We'll assume 39 bits virtual address space, allows for 512 GiB of memory (called Sv39)
- Implemented with page table indexed by Virtual Page Number (VPN) – looks up Physical Page Number (PPN)




- Page Table Entry (PTE) also stores flags in lower bits

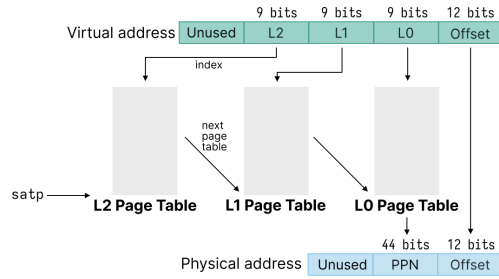
## Each Process Gets Its Own Page Table

- When you call `fork()` on a process, it will copy the page table from the product
- Turn off the write permissions so that the kernel can implement copy-on-write
- We don't need to copy the full page table – syscall `vfork()`
  - Shares all memory with the product
  - Only used in very performance sensitive programs

## 12 Page Tables

- Most programs don't use all virtual memory space, so how do we take advantage?





- Multi-level page tables save space for sparse allocations
- Given physical pages, the OS uses a free (linked) list
- Unused pages contain the next pointer in the free list
  - Physical memory gets initialized at boot, remember
- To allocate, remove from the free list, and to deallocate, add back to the free list

## 13 Page Table Implementation

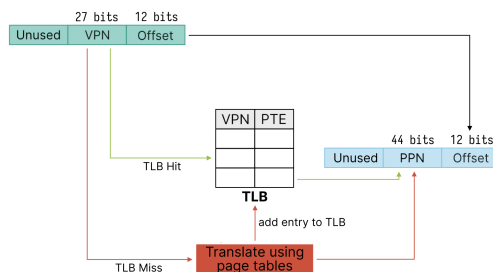
### Alignment: Memory Eventually Lines Up With Byte 0

- If pages are 4096 bytes, then pages always start when all offset bytes (12) are 0
- How many levels do I need?
  - We want each page table to fit into a single page
  - Find number of PTEs we could have in a single page  $2^{10}$

$$\text{Number of levels} = \frac{\text{Virtual bits} - \text{Offset bits}}{\text{Index bits}}$$

### Page Table for Every Memory Access is Slow


- Need to follow pointers across multiple levels of page tables
- We'll likely access the same page multiple times temporally close to the first
  - Process may only need a few VPN-PPN at a time
  - Use a computer science class – Caching!



TLB – Translation Look-aside Buffer

## Context Switching Requires Handling the TLB

- Can flush the cache or attach a PID to the TLB
- RISC-V & most implementations just flush the cache
- `sbrk` call grows or shrinks your heap, but recall, the stack has a set limit
- To grow, will grab pages from the free lists to fulfill the request
  - Kernel sets `PTE_V` (a valid bit) and other permissions
- Difficult to use in memory allocations, rarely shrinks the heap – will stay claimed as the kernel can't free pages
- Memory allocations use `mmap` to bring in large blocks of virtual memory



- The kernel can allow for processes to access fixed virtual memory/data without using a system call

## Page Faults

- Type of exception for virtual memory accesses
  - Generated if it cannot find a translation or the permission check fails
- Allows the OS to handle it
- MMU is the hardware that uses page tables:
  - May be a single page table (wasteful)
  - Use kernel-allocated pages from a free list
  - Be multi-level to save pages for sparse allocations
  - Use a TLB to speed up memory accesses

## 14 Virtual Memory Lab Primer

Not covered in these notes.

## 15 Priority Scheduling and Memory Mapping

### Dynamic Priority Scheduling

- May also be called feedback scheduling
- We let the algorithm manage priorities – we set time slices and measure CPU usage
- Increase the priority of processes that don't use their full time slices, and decrease priority of those that do
- Each process gets assigned a priority when started,  $P_n$
- Pick the lowest number to schedule, if it yields, then pick the next lowest
  - Break ties with arrival order
  - If a lower number becomes ready, switch to it
- Record how much time each process executes for in this priority interval,  $C_n$
- At the end of the priority interval, update each process

$$P_n = \frac{P_n}{2} + C_n$$

- Reset  $C_n$  back to 0 at the end of the priority interval